



VedaDB

The World's First Zero-Compromise Multi-Model Database

Bootstrapped

Pre-Seed Stage

Beta Ready

2,452 Files Built

Shubahm Mehta — Founder

shubahm.mehta@gmail.com

tiennesdm.github.io/vedadb-website

THE PROBLEM

Database Sprawl Costs \$50B+ Every Year

Modern applications need multiple data models — each requiring a separate database

7 Separate Databases

- SQL (PostgreSQL) for transactions
- Vector (Pinecone) for AI search
- Graph (Neo4j) for relationships
- Document (MongoDB) for JSON data
- Cache (Redis) for speed
- Search (Elastic) for discovery
- Time Series (InfluxDB) for metrics

The Cost

- Different APIs, query languages, and data models
- Complex ETL pipelines for cross-model queries
- Multiple backup, monitoring, and scaling systems
- 7x operational complexity
- 60% higher total cost of ownership

\$50B+

Wasted on DB Sprawl

7x

Complexity

60%

Higher TCO

3-5yr

Integration Time

THE SOLUTION

7 Engines. 1 System. Zero Compromises.

All 7 data models unified under a single query layer

SQL

ACID transactions, complex joins, window functions

Vector

HNSW indexing, cosine similarity, embeddings for AI/RAG

Graph

Cypher queries, traversal, PageRank, community detection

Document

JSON-native, flexible schema, nested queries

Cache

In-memory speed, LRU eviction, TTL support

Search

Full-text, fuzzy matching, faceted search, relevance scoring

Time Series

Columnar compression, rollups, retention policies

Key Innovation: Query across multiple engines in a single SQL statement. Join SQL tables with vector embeddings. Filter graph results with full-text search. No other database does this natively.

Built-In Intelligence

AI runs inside the database — no external APIs, no latency, no extra cost

Native AI SQL Functions (14)

EMBED()

AI_SIMILARITY()

AI_CLASSIFY()

AI_SENTIMENT()

AI_KEYWORDS()

AI_REDACT()

HYBRID_RANK()

VECTOR_DISTANCE()

RECOMMEND_VECTOR_INDEX()

VECTOR_NORM()

L2_NORMALIZE()

DP_NOISE_LAPLACE()

DP_NOISE_GAUSSIAN()

```
SELECT AI_CLASSIFY(ticket, 'urgency'), AI_SENTIMENT(body)
FROM support_tickets
WHERE created_at > NOW() - INTERVAL '1 day'
```

Key AI Capabilities

NL2SQL

Ask in English, get SQL results. Powered by Anthropic Claude with schema-aware context building.

AutoTuner

Bayesian Optimization with Gaussian Process. Automatically tunes 200+ parameters. 40% latency reduction.

Semantic Cache

LRU-based embedding deduplication. 40-80% cost reduction on repeated embedding operations.

How It Works

A unified query processor above 7 specialized storage engines

2,452

Total Files

1,593

Go Files

53

Rust SIMD Files

73

Workbench Files

LAYER	COMPONENTS	TECHNOLOGY
Query Layer	SQL Parser, Planner, Optimizer, Router, AI Optimizer	Go
7 Storage Engines	SQL, Vector, Graph, Document, Cache, Search, Time Series	Go + Rust SIMD
AI Layer	NL2SQL, AutoTuner, 14 SQL Functions, Semantic Cache	Go + Anthropic API
Interfaces	REST API, WebSocket, gRPC, 8 Language Drivers	Go / Protocol Buffers
Storage	WAL, Replication, Compression, Encryption	Go
Deployment	Docker, Kubernetes, Helm Charts	YAML / Container

What We've Built

Bootstrapped. 2,452 files. 7 engines. Beta-ready.

Completed

✓	All 7 storage engines (SQL, Vector, Graph, Document, Cache, Search, TS)
✓	Query processor with SQL parser, planner, optimizer
✓	8 language drivers (Python, Node, Go, Java, Rust, .NET, Ruby, PHP)
✓	Web workbench (React 18 + TypeScript + Tailwind)
✓	14 AI SQL functions + NL2SQL + AutoTuner
✓	Rust SIMD acceleration engine (AVX2/NEON)
✓	Docker, Kubernetes, Helm deployment
○	Final integration testing (current focus)

Roadmap

NOW	Final testing, beta release
M3	Public beta, 100+ beta users
M6	Cloud service launch
M9	First paying customers
M12	\$500K ARR target

Massive Opportunity

Database consolidation is the biggest trend in enterprise infrastructure

\$132B

DBMS Market (2024)

\$406B

Projected by 2034

16.2%

Multi-Model CAGR

\$50B+

Wasted on DB Sprawl

Market validation: IBM acquired Confluent for \$11B, MongoDB acquired Voyage AI for \$220M. Database consolidation is the defining trend of this decade.

SEGMENT	TAM	SAM	SOM	CAGR
DBMS Market	\$406B (2034)	\$32B (Multi-Model)	\$2.5B (Addressable)	23%

Freemium with Enterprise Upsell

Open source core, cloud hosting, enterprise features

TIER	PRICE	TARGET	KEY FEATURES
Open Source	Free	Developers	All 7 engines + 14 AI functions, single-node, community support
Cloud	\$299/mo	Startups	Managed hosting, AI AutoTuner, NL2SQL, backups, monitoring
Enterprise	\$999/mo	Fortune 500	Clustering, RBAC, SSO, audit, SLA, advanced AI analytics

85%

Gross Margins

18x

LTV / CAC

Land

& Expand

COMPETITION

Competitive Advantage

Native integration where others bolt on

VS.	THEIR APPROACH	VEDADB ADVANTAGE
PostgreSQL + Extensions	Separate extensions (pgvector, timescaledb) with different APIs	7 native engines, one API, one system
MongoDB Atlas	Document-only. Need separate DBs for other models.	Document + 6 more engines. True cross-model queries.
Running 7 Separate DBs	ETL pipelines, data movement, multiple ops teams	Zero data movement. One query language. One ops team.

3-6x

Faster Queries

60%

TCO Savings

7 → 1

DBs Consolidated

0

Data Movement

Execution-Focused Solo Founder

Solo founder, deep technical execution. Stripe, GitHub, Cloudflare — all started solo.

Shubahm Mehta — Founder

- Solo-built 2,452 files (1,593 Go, 53 Rust SIMD)
- Deep database internals expertise
- Query optimization, storage engines, vector operations
- Production data systems track record
- Bootstrapped from zero to beta-ready

Hiring Plan (Post Pre-Seed)

- | | |
|----|---|
| #1 | Senior Go Systems Engineer (consensus, replication) |
| #2 | DevRel Engineer (community, docs, drivers) |
| #3 | Full-Stack Engineer (cloud dashboard, billing) |

I built a database most teams of 10 couldn't ship. Pre-seed accelerates hiring and go-to-market.

THE ASK

Pre-Seed Round

Accelerating from beta to market

40%

Engineering

30%

Cloud Infrastructure

20%

Go-to-Market

10%

Operations

Milestones

M3	Complete testing, public beta release
M6	100+ beta users, cloud service launch
M9	First paying customers
M12	\$500K ARR target
M18	Break-even, Series A ready

Why Now

AI Adoption Exploding

Every app needs vector search + traditional queries. Database consolidation is the defining trend.

Market Timing

IBM bought Confluent (\$11B), MongoDB bought Voyage AI (\$220M). Multi-model is the future.

shubahm.mehta@gmail.com — tiennesdm.github.io/vedadb-website